Simulating an Interplanetary Internet

Cody Auch, Jacob Janzen, Mark Lysack

April 10, 2024

Abstract

This paper discusses interplanetary networks, the challenges that can arise from them that make them a unique challenge compared to shorter-range computer networks, and how our current network protocols can be applied in this application. In addition, we look at how one can use conventional network simulators to test interplanetary network implementations. We used the network simulator NS-3 [9] to test various transport layer protocols against each other in a network topology that we designed in the network simulator to emulate long-distance communications. Using this network simulation, we were able to compare UDP, TCP, and TCP New Reno protocols against each other to see how they each performed across such large distances with high delays. UDP preformed the best with New Reno and TCP closely paired together.

Our implementation can be found at `https://github.com/codyauch/IPN-Project`. The physics simulation was mostly designed by Mark and Cody. The network simulation in NS-3 was done by Jacob, and the analysis was done by Mark.

1 Introduction

In the past decade interest in returning back to space has increased internationally. Many organizations have started sending remote vehicles into space for the first time and there are talks of a permanent lunar base being created. With an increased amount of satellites and potentially permanent bases outside of Earths orbit, the demand on the interplanetary network (IPN) is increased. The IPN has a series of unique challenges such as a high amount of error, constantly changing network topology, and long propagation delay. The solution to these problems is implementing a delay tolerant network. While space organizations have maintained the IPN with a set of standard protocols,

these may not be sufficient for an expanding network. The simulator proposed shows how an IPN can handle different stages of development of an IPN as well as increased throughput requirements. Before these technological changes can be implemented simulations need to be created to lower the risk of failure given the high cost of failure. Often satellites will need to be in operation for years or even decades. The hardware and software in satellites can often not be updated very well. Thus it will be very relevant to investigate these issues well before they are needed.

2 Previous Work

In their 2017 Masters Thesis, Md Monjurul Islam Khan [4] from University of Manitoba wrote on how to create a satellite network simulation based on NS-3. They used a point to point NS-3 module to create a network of 2 levels of orbiting satellites and fixed points on the earth which acts as uplinks to the network. The paper analyzed TCP HighSpeed against TCP Vegas slow start. However this paper stops at expanding the network beyond earth orbit, and rather focuses on Earth orbit communication.

SNS3 is a simulation library which uses NS-3 as a basis to create a framework which to build simulations for geostationary satellites in the European continent. It does not expand to an interplanetary network [10].

For non simulation work done on the interplanetary internet, the leading authority is Consultative Committee for Space Data Systems (CCSDS) [1]. They maintain and publish the standards for interplanetary internet, including the SCPS TP, a reliable transport protocol used for delay tolerant networks when UDP or TCP are not effective. [11] Active since 1982 and managed in collaboration with several national and international space organizations. A contemporary simulation of a current interplanetary network or earth orbit network would follow the protocols and standards laid out by CCSDS.

Several public GitHub repositories with satellite simulators can be found, for instance:

- `https://github.com/szymonwieloch/DTN/tree/master`

- `https://goodingc.github.io/ipn_sim/web-app/`

Both of these limit themselves to earth orbit simulations.

3　Interplanetary Internet and Background

An interplanetary network is simply a network which where the nodes are distributed across the surface or in the orbits of different planets or other celestial objects. In the year of 2024 the current interplanetary is limited to Mars, Earth, Earth orbit, Moon orbit, and satellites orbiting the sun. The interplanetary network is also expanded to deep space satellites, like the Voyager probes. A sample message path across the network could be that a message is passed from a large network antenna on earth to a earth orbiting satellite, then when it has direct line of sight with MAVEN (a satellite orbiting mars) it sends the message. MAVEN then relays the message to a Mars rover like Odyssey. This description hides the complexity which arises in an interplanetary network which may seem like a extremely simple network.

In a traditional network a majority of the network traffic will occur over a wired connection, especially for transfers of large amounts of data. This means that the error is relatively low. This is no longer the case in the case of satellites sending messages over potentially millions of kilometers. Atmosphere in the Earth orbit, solar radiation, and dust cause increased error in sending the message between entities. Inverse square law dictates that over these large periods the energy needed to send the message so that it can be read at its destination. Even the Doppler effect is needed to be accounted for to tune into radio frequencies on the ISS. [2] High error means that sending a message can no longer be assumed to be sent reliably and has to be compensated for. Beyond high error, large distances between entities and the physical limitation that messages cannot travel faster than the speed of light. Light, or any other energy takes approximately 8 minutes to get from the Sun to Earth, [7] and thus for a round trip will take 16 minutes to complete a simple handshake, given that the messages were sent correctly.

Other than physical limitations a interplanetary network has a rapidly changing topology at anytime. For instance satellites which orbit planets can have their field of view blocked as they move over the horizon. While terrestrial networks also have rapidly changing networks at the edges, but the difference is that the core of the network is often stable. Where depending on the hour, day, year, the core of the interplanetary network can change drastically. To be able to connect to the network there are no well known places to access the network, except for Earth.

To combat these challenges several protocols have been created. The SCPS TP is a transport protocol which is used on the network layer. UDP is used in the lowest reliable connections between nodes, and TCP is used in the most reliable network conditions. When a reliable and efficient method of network transportation is needed, then SCPS TP is used. It is a modified

Figure 1: A sample of the current interplanetary network

version of TCP which uses a slow start protocol, reduced handshakes, reliable data transfer with reduced confirmations. Other protocols like store and load, and Licklider are used to also overcome these problems. In this paper since no complete implementation of SCPS TP has been created for NS-3, New Reno TCP was used instead.

The rest of the paper remains to simulate these unique challenges, as well as see how protocols operate at these different situations.

4   Methodology

4.1   Physics Engine

A significant challenge in modeling a IPN is modeling the rapidly changing topology. To simulate these rapid changing conditions a physics engine was created to model the orbits of several entities in the solar system. Several assumptions were made in the simulator for the sake of simplicity as the goal of the simulator is not to create a perfect model of the solar system and the entities inside of the solar system, but rather conditions which a IPN would exist in.

The first assumption made was that all orbits are circular. Realistically, orbits are oval shaped, with two focal points. A result of this assumption is the elimination of any difference in orbital speed due to Kepler's second law.

Normally a planet will accelerate or decelerate depending on where it is in its orbit around the sun. The orbital radius was chosen by taking the average orbital radius of the real life celestial bodies. For example Earth's average orbital radius is 149.6 billion meters [7], and therefore orbits around the sun in a circular orbit of radius 149.6 billion meters in our model. Another result of this assumption is that orbits are considered to be centered around the body which it orbits. The Mars orbit has a focal point that is further in space, making it more oval shaped than the earth orbit. [8] While this could make changes to a simulated network in some very specific situations, the primary behaviour of a solar system is maintained, like retrograde orbits, periodical alignment, predictable behaviour, and recursive orbits of sub orbiting object.

All entities are assumed to have a circular shape, with a constant radius. Due to the scale of the objects, both large and small, this has a very minor effect.

The whole physics simulation has been done on a 2D plane. While planetary orbits were not perfectly aligned, on the astronomical scale which we are working on, assuming that all their orbits are contained on the same plane will not effect the simulation significantly. Where this could be an issue is for when there are more than two satellites orbiting the same entity, with a relatively low orbit. For example satellites orbiting Earth. While there are some complex satellite interactions which are removed it does still simulate satellites having to form a network around the entity it is orbiting, it just requires less satellites then if it needed to form a 3-D network.

Messages between entities were assumed to be sent at the speed of light in a vacuum, as all forms of radio, lasers, or other wireless messages are forms of electromagnetic radiation, and thus travel at the speed of light. Since the propagation delay in an IPN is much higher than in terrestrial networks, other delays such as processing delays were ignored as being negligible.

Based on these simplifications, a physical simulation with several goals. To calculate if a line of sight could be made between two communicating entities, that is if a message can be sent, to calculate the error rate when sending a message, and to calculate the message propagation delay for the message to be sent. Line of sight was achieved using simple geometry calculations. All entities are able to block signals if the ray passes through the radius of the entity. One of the largest factors contributing to transmission error is the distance traveled by the electromagnetic signal between transmitting and receiving nodes. By generalizing Friis' transmission formula[3] for isotropic emitters and receivers, we can obtain the ratio of signal intensity received to signal intensity transmitted, known as Free Space Path Loss (FSPL). We can use FSPL as our transmission error rate, giving the following expression:

$ErrorRate = \left(\frac{4\pi d}{\lambda}\right)^2$ where d is the distance between nodes in metres and λ is the wavelength of the electromagnetic wave used in transmission. As much of space communications utilizes the Ka-band of frequencies (27–40 GHz)[5] we opted to use a transmission frequency of 30Ghz. The transmission wavelength can be then calculated according to $\lambda = \frac{c}{f}$, where c is the speed of light and f is the transmission frequency, giving a transmission wavelength of about 9.99 mm. Propagation delay is a simple function of distance and the speed of light.


Figure 2: A visual output of the physics model. Left: Mars and Earth orbiting the sun. Right: The ISS (overlapping the center) and the Moon orbiting the Earth

The physics simulation modeled the Sun, Earth, Moon, International Space Station, Mars, and the Mars Orbiter. All were able to connect to the network except for the the Sun and the Moon. The input of the physics engine is a time t from some preset initial conditions which outputs the state of all entities. Which is fed into the NS-3 simulation.

4.2   Network Simulation

We used NS-3 [9] to implement a simulation of this network. Our final design consisted of a set of routers representing the different entities in our physics simulation which were each linked to every other router using a point-to-point channel. Essentially, this allowed us to have a wired network with configurable delay and error along with the ability to take certain channels down if they are obstructed that we could make behave like an interplanetary network would without having to deal with the complexities of configuring wireless communication. In addition to the mesh of routers, we added two individual nodes to the network with zero delay that can be connected to any one router in the network. One of these two nodes acts as a sender and the other one acts as a receiver. This configuration of sender and receiver as separate nodes

made it easier to decouple the applications from the network topology and allowed us to programmatically create arbitrary topologies.

NS-3 is a useful library, but we did run into a number of problems over the course of our implementation of our simulation. The first problem came when we first tried to compile the library itself and ran into a bug in the build script that caused all executables to fail to run if there was the word "scratch" anywhere in the path due to a specific subdirectory of NS-3 which also happens to be called "scratch". When first testing it out, one of us had put NS-3 into a directory that happened to contain that word. We made a pull request to fix this issue and it was merged by the NS-3 maintainers which allowed us to continue.

The next issue we ran into was that we had decided to use the Python bindings to allow us to easily use Python for the physics simulation. This was a mistake. Although NS-3 does have Python bindings that do make it easier to run it as a script after writing, there is no documentation at all for the Python bindings and some basic features such as setting the error rate on a channel or scheduling a task to run during the simulation have no Python support and must use inline C++. The lack of support for task scheduling in Python is particularly bad as it forced us to schedule a C++ function which itself called Python code. It also forced us to use global state as `CPyCpyy` does not have support for arbitrary function parameters when calling Python code as far as we could tell (this feature is not documented and we only figured out about it because the few Python examples in the NS-3 source code that do use scheduling also call Python from C++ with this method).

Before we settled on using point-to-point to emulate a wireless network, we attempted to use the existing wireless networking features of NS-3. These features are very powerful and even have a positioning and velocity system for nodes in the network which would be very useful. We ran into a lot of issues with this system though. Firstly, there only seems to be real support for Wi-Fi which is not particularly useful for long distance unlike radio. This meant that we would have to greatly decrease the distances that we were simulating. This is fine and would still allow a perfectly reasonable simulation so long as we increased the distances again in our computations while analyzing the data. The other problem that we ran into is the fact that the wireless networking would quickly require us to go well outside the scope of our project. Determining the gain and energy required to configure our satellites to send and receive messages correctly is more of an engineering problem that we were not interested in solving.

After settling on using point-to-point, UDP was easy to get working but TCP proved to be very challenging. Our initial implementation of our topology seemed to have an issue where messages could be routed in one direction but

not the other. This meant that when TCP tried to establish a connection it would always fail. It still is not entirely clear what was configured incorrectly, but after changing how we were creating channels, TCP was finally working. We ran into more issues with TCP though, when we started adding delay to the channels. It turns out that NS-3 sockets have a rather low time-to-live and there do not seem to be any global options for configuring sockets. We were left with two options: implement TCP from scratch using custom-built sockets for each node in the network or ensure that the messages that we do send have lower delay than the TTL for our sockets. We opted for the latter. To make sure that TCP messages could successfully ACK, we found a value that we could divide all times by to make sure that almost every message that could feasibly be sent would be received before its TTL expired. This does not change any distances or error rates and effectively only changes the unit we use for time from s to $\frac{s}{26}$ which has an easy conversion back to s. Finally, we implemented a version of our program using the New Reno variant of TCP. Thankfully, NS-3 made this easy for us. It worked perfectly after assigning a configuration value before setting up our communication channels.

During the work on this project we contributed to the open source NS-3 project by committing a change to fix a bug which occurred when running the program on different operating systems.

5 Data

The NS-3 simulation ran on the basis of transferring data between Earth and Mars.

Data was scraped from the trace output that was generated while running the ns3 simulation. Trace data was given when a message was queued, dequeued, and received at the destination. This data also included the amount of data being sent in bytes. There is other data on the messages being sent, but was not tracked for the purpose of this project. The trace data was parsed using regex expressions to extract the needed numbers from the string. This data was then fed into a pandas dataframe. A unique id had to be generated for each message since once a message was received the id was reused. This was done by tracking when each message was received with id x and after incrementing a number by one, say y. The unique id was then created by hashing xy. Once the data frame was created it had to be merged on itself to have one row of the dataframe represent exactly one message rather then one row be either queue, dequeue, or received.

Time from being queued to received, time being queued, and message time in transit were all calculated based off of the data.

$$\text{total\_package\_time} = \text{time\_recieved} - \text{time\_queued}$$

$$\text{time\_queued} = \text{time\_dequeued} - \text{time\_queued}$$

$$\text{total\_time} = \text{time\_received} - \text{time\_dequeued}$$

After the data was processed it was saved in a csv file with one row representing exactly one message that was sent over the network. This file was close to 2.5 GB in size.

Statistics were calculated on the entire column of the calculated rows, and bytes sent. Messages were also grouped together which occurred during a similar time window and further averages were calculated to find how the average changed over time of the simulation. These calculations are shown in the graphs in the analysis section.

6   Analysis

6.1   UDP

Figure 3: UDP analyzed

UDP had an average time from being queued to received of 283.64 seconds, an average time being queued of 21.92 seconds, and an average message time in transit 261.72 seconds. The average amount of data sent was 1052.0 bytes. Notice the scientific notation in the corner of the graph.

We believe that UDP sends a large number of messages at the start that are all queued, processed, and sent. After those messages are sent, the next batch is prepared immediately. This curve could be caused by the queuing delay from the large number of messages that are sent. UDP does not wait for messages to be received and will therefore create a large backlog which slows everything down when dequeuing messages. We see gap between lines which seems to be time required to send the messages to the planet from an orbiting body.

6.2   TCP

Figure 4: TCP analyzed

The average amount of bytes transferred is 320.44 bytes. This is likely from the difference in the amounts of bytes transferred in an ACK message compared to when data is being sent.

Unlike in UDP the motion of the planets is not observable in the graph. However this is because there are several different paths which the TCP message could have taken which can be seen. Since TCP requires a higher amount of base time to establish a connection compared to UDP, this means that when it "hops" between entities it adds a larger base time.

At the start of the time we can see 4 separate paths. This would be the different times that it would take for any of the 4 paths between Earth and Mars, that is Earth to the ISS to the Mars Orbiter then to Mars, Earth to the Mars Orbiter to Mars, Earth to the ISS to Mars, and finally Earth to Mars. After one of the entities is blocking a direct signal it becomes only that 3 of these options possible, with likely only two actual paths being viable.

With less available bandwidth it becomes slower on average. Smaller time may also be correlated to the minority of messages which were sent with a smaller message size.


Figure 5: TCP in the future analyzed

The second set of graphs show TCP in some time in the future. We can see that the time that messages take to transfer does in fact change over time as the entities change distances. We can also see a similar pattern of behaviour changing as some topology changes in the model.

6.3   New Reno

New Reno, a modified version of TCP used to help reduce congestion [6], was used in place of SCP TP. New Reno behaved very similar to TCP. Its average data transfer was 307.18 bytes.

It preformed marginally better in its time from being queued to the message received. This was mostly due to a decrease in time that it was in queue, which is to be expected. There was a very marginal increase in time in transit however.

Similar to TCP we can see the two stages which occur in the simulation. One where most of the messages were likely being routed from directly from Earth to Mars, which then gets transformed into having to route through one of the external satellites orbiting each entity, which decreases the average time it takes to transmit the message.
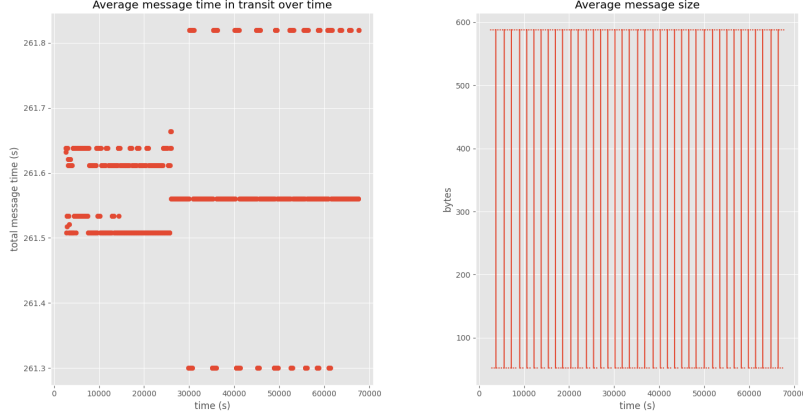
11

Figure 6: TCP in the future analyzed

6.4   Protocol Comparison

Since a large amount of propagation delay dominates the amount of time that it takes for messages total transmit time from being queued to being received at the appropriate node, other increases in performance mean only a slight decrease in time sending the message. An important statistic to pay attention to is that UDP almost always sent a message with a size of 1052 bytes. This means that even though TCP, New Reno, and UDP had sent a similar number of messages, the actual throughput for UDP was significantly higher, which is much more significant in an IPN situation. In the data graphs we can see that not only was the maximum data being sent was less than UDP, but also had to send messages of very little size, most likely acknowledgments, as well. New Reno was seen to improve on the queuing time compared to both TCP and UDP. Which would mean in a high reliability situation and in situations where higher traffic is sent through a network with a relatively low bandwidth, like a IPN, where satellites are limited by how much data they can transmit, New Reno's effect can be more significant. This should only be used in a connection where propagation delay is less of an issue like between planet and satellite communication. Queuing time in UDP took a more significant percentage of its time than both TCP and New Reno. In general however, UDP is the best option for primary communication between Mars and Earth communication.

12

7 Future Work

Future work that could be done would be to implement or partially implement SCP TP as a network protocol which could be used and compared to the protocols analyzed in this paper. Other protocols like the Linklider or load and store protocols could also be implemented, to create a more realistic IPN.

The physics engine could be expanded to included 3-D orbits, particularly around planets, where networks of satellites could be created to create a proper level of communication that is maintained today, with geocentric and polar orbiting satellites. Connection points from a planet to off the planet could have a direction which it needs to send a message in for it to be able to connect. For example a rover on Mars in our simulation could technically send a message in any direction, even if they were sending it through the surface of the planet.The error model in the physics engine could be expanded to include interference like solar radiation, atmosphere, and other interference.

Other stats could be extracted from the NS-3 model such as packet loss, message path, handshake time. It could expand so that the network was not just focused on Earth to Mars communication. It could be expanded to include other celestial objects in the solar system, as well as potential future objects.

The NS-3 simulation could dynamically choose the network protocol to use based on the level of error, message size, estimated propagation error, and other factors. This could build a model which uses not only one protocol but all of them and leverages the advantages of each when it is necessary. This combined with a more sophisticated model of the solar system would create a very competent simulation of the current operation of the IPN, and could then be expanded or tweaked to find if expansions to it could be handled properly. This would need a combination of all of the suggestions above to be successful.

8 Conclusion

In this paper we outlined how we created a model of an interplanetary internet (IPN) using a physics engine which fed into a NS-3 simulation. The physics simulation used a simplified model of orbiting planets and satellites as well as physics calculations to create data which would be the output for any input time t. The NS-3 model used a point to point base which programmatically took the data from the physics engine and created nodes which acted as routers. A simulation was ran which transferred data between Earth and Mars, using satellites and other entities to relay messages around blocked

connections. Analysis was done on three different network protocols used in the IPN, that is UDP, TCP and New Reno. UDP was found to have the most throughput out of all of the three protocols, while New Reno reduced the amount of time a message was queued.

References

[1] CCSDS. *Implementation Software and Guides*. URL: `https://public.ccsds.org/about/charter.aspx`.

[2] *Compensating for Doppler When Contacting the ISS (International Space Station) — qsl.net*. `https://www.qsl.net/ah6rh/am-radio/spacecomm/doppler-and-the-iss.html`. [Accessed 10-04-2024].

[3] H.T. Friis. "A Note on a Simple Transmission Formula". In: *Proceedings of the IRE* 34.5 (1946), pp. 254–256. DOI: `10.1109/JRPROC.1946.234568`.

[4] Md Monjurul Islam Khan. *Development of an ns-3 Based Network Simulator for Space Telemetry*. 2017. URL: `https://mspace.lib.umanitoba.ca/server/api/core/bitstreams/b1315298-9ccd-43bd-ade5-3064ac410976/content`.

[5] D. Morabito and R. Hastrup. "Communicating with Mars during periods of solar conjunction". In: *Proceedings, IEEE Aerospace Conference*. Vol. 3. 2002, pp. 3–3. DOI: `10.1109/AERO.2002.1035260`.

[6] Lutfun Nahar et al. "Throughput analysis for TCP NewReno". In: Sept. 2016, pp. 1–6. DOI: `10.1109/CEEICT.2016.7873064`.

[7] NASA. *Earth Fact Sheet*. URL: `https://nssdc.gsfc.nasa.gov/planetary/factsheet/earthfact.html`.

[8] NASA. *Mars Fact Sheet*. URL: `https://nssdc.gsfc.nasa.gov/planetary/factsheet/marsfact.html`.

[9] nsnam. *NS-3*. Version 3.41. Feb. 10, 2024. URL: `https://www.nsnam.org/`.

[10] Jani Puttonen et al. "Satellite Model for Network Simulator 3". In: ICST, Aug. 2014. DOI: `10.4108/icst.simutools.2014.254631`.

[11] Keith Scott and Leigh Torgerson. *SCPS - TP : a satellite - enhanced TCP*. Version V1. 2004. DOI: `2014/40565`. URL: `https://hdl.handle.net/2014/40565`.